

Here we are given a data set diamonds.csv in which we have to predict the diamond price.

In [727]:

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import math
import random
import os
```

To deal with shuffling.

In [728]:

```
def seed():
    np.random.seed(42)
    random.seed(42)
    os.environ["PYTHONHASHSEED"] = "42"

seed()
```

In [729]:

```
ds = pd.read_csv('/content/diamonds.csv')
```

In [730]:

```
ds.loc[:20,:]
```

Out[730]:

	Unnamed: 0	carat	cut	color	clarity	depth	table	price	x	y	z
0	1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
5	6	0.24	Very Good	J	VVS2	62.8	57.0	336	3.94	3.96	2.48
6	7	0.24	Very Good	I	VVS1	62.3	57.0	336	3.95	3.98	2.47
7	8	0.26	Very Good	H	SI1	61.9	55.0	337	4.07	4.11	2.53
8	9	0.22	Fair	E	VS2	65.1	61.0	337	3.87	3.78	2.49
9	10	0.23	Very Good	H	VS1	59.4	61.0	338	4.00	4.05	2.39
10	11	0.30	Good	J	SI1	64.0	55.0	339	4.25	4.28	2.73
11	12	0.23	Ideal	J	VS1	62.8	56.0	340	3.93	3.90	2.46
12	13	0.22	Premium	F	SI1	60.4	61.0	342	3.88	3.84	2.33
13	14	0.31	Ideal	J	SI2	62.2	54.0	344	4.35	4.37	2.71
14	15	0.20	Premium	E	SI2	60.2	62.0	345	3.79	3.75	2.27
15	16	0.32	Premium	E	I1	60.9	58.0	345	4.38	4.42	2.68
16	17	0.30	Ideal	I	SI2	62.0	54.0	348	4.31	4.34	2.68
17	18	0.30	Good	J	SI1	63.4	54.0	351	4.23	4.29	2.70
18	19	0.30	Good	J	SI1	63.8	56.0	351	4.23	4.26	2.71

19	Unnamed: 0	20	carat	21	0.30	Very Good	cut	color	clarity	depth	table	price	4.26	4.30	2.71
20		21	0.30			Good	I	S12	63.3	56.0	351	4.26	4.30	2.71	

In [731]:

```
ds.describe()
```

Out[731]:

	Unnamed: 0	carat	depth	table	price	x	y	z
count	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000	53940.000000
mean	26970.500000	0.797940	61.749405	57.457184	3932.799722	5.731157	5.734526	3.538734
std	15571.281097	0.474011	1.432621	2.234491	3989.439738	1.121761	1.142135	0.705699
min	1.000000	0.200000	43.000000	43.000000	326.000000	0.000000	0.000000	0.000000
25%	13485.750000	0.400000	61.000000	56.000000	950.000000	4.710000	4.720000	2.910000
50%	26970.500000	0.700000	61.800000	57.000000	2401.000000	5.700000	5.710000	3.530000
75%	40455.250000	1.040000	62.500000	59.000000	5324.250000	6.540000	6.540000	4.040000
max	53940.000000	5.010000	79.000000	95.000000	18823.000000	10.740000	58.900000	31.800000

Null Values

Checking for the null values

In [732]:

```
ds.isna().sum()
```

Out[732]:

```
Unnamed: 0      0
carat           0
cut             0
color           0
clarity         0
depth           0
table           0
price           0
x               0
y               0
z               0
dtype: int64
```

Checking unique Values

In [733]:

```
ds.nunique()
```

Out[733]:

```
Unnamed: 0      53940
carat           273
cut              5
color            7
clarity          8
depth           184
table           127
price          11602
x               554
y               552
z               375
dtype: int64
```

We can drop the Unnamed Column as it is just indexing

In [734]:

```
ds.drop(axis="columns", labels="Unnamed: 0", inplace=True)
ds.loc[:10,:]
```

Out[734]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75
5	0.24	Very Good	J	VVS2	62.8	57.0	336	3.94	3.96	2.48
6	0.24	Very Good	I	VVS1	62.3	57.0	336	3.95	3.98	2.47
7	0.26	Very Good	H	SI1	61.9	55.0	337	4.07	4.11	2.53
8	0.22	Fair	E	VS2	65.1	61.0	337	3.87	3.78	2.49
9	0.23	Very Good	H	VS1	59.4	61.0	338	4.00	4.05	2.39
10	0.30	Good	J	SI1	64.0	55.0	339	4.25	4.28	2.73

We see x(length), y(width) and z(depth) has minimum value 0 which indicates these are wrong values so we can remove those rows

In [735]:

```
ds = ds.drop(ds[ds["x"]==0].index)
ds = ds.drop(ds[ds["y"]==0].index)
ds = ds.drop(ds[ds["z"]==0].index)
```

In [736]:

```
ds.describe()
```

Out[736]:

	carat	depth	table	price	x	y	z
count	53920.000000	53920.000000	53920.000000	53920.000000	53920.000000	53920.000000	53920.000000
mean	0.797698	61.749514	57.456834	3930.993231	5.731627	5.734887	3.540046
std	0.473795	1.432331	2.234064	3987.280446	1.119423	1.140126	0.702530
min	0.200000	43.000000	43.000000	326.000000	3.730000	3.680000	1.070000
25%	0.400000	61.000000	56.000000	949.000000	4.710000	4.720000	2.910000
50%	0.700000	61.800000	57.000000	2401.000000	5.700000	5.710000	3.530000
75%	1.040000	62.500000	59.000000	5323.250000	6.540000	6.540000	4.040000
max	5.010000	79.000000	95.000000	18823.000000	10.740000	58.900000	31.800000

In [737]:

```
ds.nunique()
```

Out[737]:

```
carat      273
cut         5
color       7
clarity     8
```

```
depth      184
table      127
price     11597
x          553
y          550
z          374
dtype: int64
```

Now we see cut, color and clarity has 5, 7 and 8 unique values. Lets check these.

```
In [738]:
```

```
ds["cut"].unique()
```

```
Out[738]:
```

```
array(['Ideal', 'Premium', 'Good', 'Very Good', 'Fair'], dtype=object)
```

```
In [739]:
```

```
ds["color"].unique()
```

```
Out[739]:
```

```
array(['E', 'I', 'J', 'H', 'F', 'G', 'D'], dtype=object)
```

```
In [740]:
```

```
ds["clarity"].unique()
```

```
Out[740]:
```

```
array(['SI2', 'SI1', 'VS1', 'VS2', 'VVS2', 'VVS1', 'I1', 'IF'],
      dtype=object)
```

Label Encoding

We can do label encoding on these columns. We have to take care of some cases like:

Color: J= Worst(1) D= Best(7)

Clarity: (I1 (1), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (8))

Cut: (Fair(1), Good, Very Good, Premium, Ideal(5))

```
In [741]:
```

```
ds["color"] = ds["color"].map({"J": 1, "I": 2, "H": 3, "G": 4, "F": 5, "E": 6, "D": 7})
ds["clarity"] = ds["clarity"].map({"I1": 1, "SI2": 2, "SI1": 3, "VS2": 4, "VS1": 5, "VVS2": 6, "VVS1": 7, "IF": 8})
ds["cut"] = ds["cut"].map({"Fair": 1, "Good": 2, "Very Good": 3, "Premium": 4, "Ideal": 5})
```

```
In [742]:
```

```
ds.loc[:10,:]
```

```
Out[742]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	4	6	3	59.8	61.0	326	3.89	3.84	2.31
2	0.23	2	6	5	56.9	65.0	327	4.05	4.07	2.31
3	0.29	4	2	4	62.4	58.0	334	4.20	4.23	2.63
4	0.31	2	1	2	63.3	58.0	335	4.34	4.35	2.75
5	0.24	3	1	6	62.8	57.0	336	3.94	3.96	2.48
6	0.24	3	2	7	62.3	57.0	336	3.95	3.98	2.47

	carat	cut	color	clarity	depth	table	price	x	y	z
7	0.26	3	3	3	61.9	55.0	337	4.07	4.11	2.53
8	0.22	1	6	4	65.1	61.0	337	3.87	3.78	2.49
9	0.23	3	3	5	59.4	61.0	338	4.00	4.05	2.39
10	0.30	2	1	3	64.0	55.0	339	4.25	4.28	2.73

Duplicate Rows

We are removing the duplicate rows if present.

In [743]:

```
ds.duplicated().sum()
```

Out[743]:

145

In [744]:

```
ds.drop(axis="rows", labels=ds.index[ds.duplicated()], inplace=True)
ds.duplicated().sum()
```

Out[744]:

0

Feature Scaling

We are now doing the feature scaling for columns x, y, z, table and depth using min-max normalization

In [745]:

```
x = ds["x"] - np.min(ds["x"])
y = np.max(ds["x"]) - np.min(ds["x"])
ds["x"] = x/y
ds.loc[:20,:]
```

Out[745]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	61.5	55.0	326	0.031384	3.98	2.43
1	0.21	4	6	3	59.8	61.0	326	0.022825	3.84	2.31
2	0.23	2	6	5	56.9	65.0	327	0.045649	4.07	2.31
3	0.29	4	2	4	62.4	58.0	334	0.067047	4.23	2.63
4	0.31	2	1	2	63.3	58.0	335	0.087019	4.35	2.75
5	0.24	3	1	6	62.8	57.0	336	0.029957	3.96	2.48
6	0.24	3	2	7	62.3	57.0	336	0.031384	3.98	2.47
7	0.26	3	3	3	61.9	55.0	337	0.048502	4.11	2.53
8	0.22	1	6	4	65.1	61.0	337	0.019971	3.78	2.49
9	0.23	3	3	5	59.4	61.0	338	0.038516	4.05	2.39
10	0.30	2	1	3	64.0	55.0	339	0.074180	4.28	2.73
11	0.23	5	1	5	62.8	56.0	340	0.028531	3.90	2.46
12	0.22	4	5	3	60.4	61.0	342	0.021398	3.84	2.33
13	0.31	5	1	2	62.2	54.0	344	0.088445	4.37	2.71
14	0.20	4	6	2	60.2	62.0	345	0.008559	3.75	2.27
15	0.32	4	6	1	60.9	58.0	345	0.092725	4.42	2.68

	carat	cut	color	clarity	depth	table	price	x	y	z
16	0.30	5	2	2	62.0	54.0	348	0.082739	4.34	2.68
17	0.30	2	1	3	63.4	54.0	351	0.071327	4.29	2.70
18	0.30	2	1	3	63.8	56.0	351	0.071327	4.26	2.71
19	0.30	3	1	3	62.7	59.0	351	0.068474	4.27	2.66
20	0.30	2	2	2	63.3	56.0	351	0.075606	4.30	2.71

In [746]:

```
x = ds["price"] - np.min(ds["price"])
y = np.max(ds["price"]) - np.min(ds["price"])
ds["price"] = x/y
ds.loc[:20,:]
```

Out[746]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	61.5	55.0	0.000000	0.031384	3.98	2.43
1	0.21	4	6	3	59.8	61.0	0.000000	0.022825	3.84	2.31
2	0.23	2	6	5	56.9	65.0	0.000054	0.045649	4.07	2.31
3	0.29	4	2	4	62.4	58.0	0.000433	0.067047	4.23	2.63
4	0.31	2	1	2	63.3	58.0	0.000487	0.087019	4.35	2.75
5	0.24	3	1	6	62.8	57.0	0.000541	0.029957	3.96	2.48
6	0.24	3	2	7	62.3	57.0	0.000541	0.031384	3.98	2.47
7	0.26	3	3	3	61.9	55.0	0.000595	0.048502	4.11	2.53
8	0.22	1	6	4	65.1	61.0	0.000595	0.019971	3.78	2.49
9	0.23	3	3	5	59.4	61.0	0.000649	0.038516	4.05	2.39
10	0.30	2	1	3	64.0	55.0	0.000703	0.074180	4.28	2.73
11	0.23	5	1	5	62.8	56.0	0.000757	0.028531	3.90	2.46
12	0.22	4	5	3	60.4	61.0	0.000865	0.021398	3.84	2.33
13	0.31	5	1	2	62.2	54.0	0.000973	0.088445	4.37	2.71
14	0.20	4	6	2	60.2	62.0	0.001027	0.008559	3.75	2.27
15	0.32	4	6	1	60.9	58.0	0.001027	0.092725	4.42	2.68
16	0.30	5	2	2	62.0	54.0	0.001189	0.082739	4.34	2.68
17	0.30	2	1	3	63.4	54.0	0.001352	0.071327	4.29	2.70
18	0.30	2	1	3	63.8	56.0	0.001352	0.071327	4.26	2.71
19	0.30	3	1	3	62.7	59.0	0.001352	0.068474	4.27	2.66
20	0.30	2	2	2	63.3	56.0	0.001352	0.075606	4.30	2.71

In [747]:

```
x = ds["y"] - np.min(ds["y"])
y = np.max(ds["y"]) - np.min(ds["y"])
ds["y"] = x/y
ds.loc[:10,:]
```

Out[747]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	61.5	55.0	0.000000	0.031384	0.005433	2.43
1	0.21	4	6	3	59.8	61.0	0.000000	0.022825	0.002898	2.31
2	0.23	2	6	5	56.9	65.0	0.000054	0.045649	0.007063	2.31
3	0.29	4	2	4	62.4	58.0	0.000433	0.067047	0.009960	2.63

4	0.31	2	1	2	63.3	58.0	0.000487	0.087019	0.012133	2.75
5	0.24	3	1	6	62.8	57.0	0.000541	0.029957	0.005071	2.48
6	0.24	3	2	7	62.3	57.0	0.000541	0.031384	0.005433	2.47
7	0.26	3	3	3	61.9	55.0	0.000595	0.048502	0.007787	2.53
8	0.22	1	6	4	65.1	61.0	0.000595	0.019971	0.001811	2.49
9	0.23	3	3	5	59.4	61.0	0.000649	0.038516	0.006700	2.39
10	0.30	2	1	3	64.0	55.0	0.000703	0.074180	0.010866	2.73

In [748]:

```
x = ds["z"] - np.min(ds["z"])
y = np.max(ds["z"]) - np.min(ds["z"])
ds["z"] = x/y
ds.loc[:5,:]
```

Out[748]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	61.5	55.0	0.000000	0.031384	0.005433	0.044256
1	0.21	4	6	3	59.8	61.0	0.000000	0.022825	0.002898	0.040351
2	0.23	2	6	5	56.9	65.0	0.000054	0.045649	0.007063	0.040351
3	0.29	4	2	4	62.4	58.0	0.000433	0.067047	0.009960	0.050765
4	0.31	2	1	2	63.3	58.0	0.000487	0.087019	0.012133	0.054670
5	0.24	3	1	6	62.8	57.0	0.000541	0.029957	0.005071	0.045884

In [749]:

```
x = ds["depth"] - np.min(ds["depth"])
y = np.max(ds["depth"]) - np.min(ds["depth"])
ds["depth"] = x/y
ds.loc[:10,:]
```

Out[749]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	0.513889	55.0	0.000000	0.031384	0.005433	0.044256
1	0.21	4	6	3	0.466667	61.0	0.000000	0.022825	0.002898	0.040351
2	0.23	2	6	5	0.386111	65.0	0.000054	0.045649	0.007063	0.040351
3	0.29	4	2	4	0.538889	58.0	0.000433	0.067047	0.009960	0.050765
4	0.31	2	1	2	0.563889	58.0	0.000487	0.087019	0.012133	0.054670
5	0.24	3	1	6	0.550000	57.0	0.000541	0.029957	0.005071	0.045884
6	0.24	3	2	7	0.536111	57.0	0.000541	0.031384	0.005433	0.045558
7	0.26	3	3	3	0.525000	55.0	0.000595	0.048502	0.007787	0.047511
8	0.22	1	6	4	0.613889	61.0	0.000595	0.019971	0.001811	0.046209
9	0.23	3	3	5	0.455556	61.0	0.000649	0.038516	0.006700	0.042955
10	0.30	2	1	3	0.583333	55.0	0.000703	0.074180	0.010866	0.054019

In [750]:

```
x = ds["table"] - np.min(ds["table"])
y = np.max(ds["table"]) - np.min(ds["table"])
ds["table"] = x/y
ds.loc[:10,:]
```

Out[750]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	5	6	2	0.513889	0.230769	0.000000	0.031384	0.005433	0.044256
1	0.21	4	6	3	0.466667	0.346154	0.000000	0.022825	0.002898	0.040351
2	0.23	2	6	5	0.386111	0.423077	0.000054	0.045649	0.007063	0.040351
3	0.29	4	2	4	0.538889	0.288462	0.000433	0.067047	0.009960	0.050765
4	0.31	2	1	2	0.563889	0.288462	0.000487	0.087019	0.012133	0.054670
5	0.24	3	1	6	0.550000	0.269231	0.000541	0.029957	0.005071	0.045884
6	0.24	3	2	7	0.536111	0.269231	0.000541	0.031384	0.005433	0.045558
7	0.26	3	3	3	0.525000	0.230769	0.000595	0.048502	0.007787	0.047511
8	0.22	1	6	4	0.613889	0.346154	0.000595	0.019971	0.001811	0.046209
9	0.23	3	3	5	0.455556	0.346154	0.000649	0.038516	0.006700	0.042955
10	0.30	2	1	3	0.583333	0.230769	0.000703	0.074180	0.010866	0.054019

In [751]:

```
ds.describe()
```

Out[751]:

	carat	cut	color	clarity	depth	table	price	x
count	53775.000000	53775.000000	53775.000000	53775.000000	53775.000000	53775.000000	53775.000000	53775.000000
mean	0.797536	3.904231	4.406267	4.052366	0.520784	0.278035	0.194908	0.285532
std	0.473169	1.116097	1.701271	1.646733	0.039712	0.042947	0.215490	0.159574
min	0.200000	1.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000
25%	0.400000	3.000000	3.000000	3.000000	0.500000	0.250000	0.033789	0.139800
50%	0.700000	4.000000	4.000000	4.000000	0.522222	0.269231	0.112180	0.281027
75%	1.040000	5.000000	6.000000	5.000000	0.541667	0.307692	0.270206	0.400856
max	5.010000	5.000000	7.000000	8.000000	1.000000	1.000000	1.000000	1.000000


Correlation


In [752]:

```
plt.figure(figsize=(11,11))
sns.heatmap(ds.corr(), annot=True)
```

Out[752]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f8950970050>





"x", "y" and "z" show a high correlation to the "price" column. So we can drop any two columns. We also see "carat" is very much dependent on "x", "y" and "z" as it is obvious if length, width and depth increases carat will increase. So let's drop carat as it has the highest correlation.

"depth", "cut" and "table" show low correlation.

In [753]:

```
ds.drop(axis="columns", labels="carat", inplace=True)
```

In [754]:

```
ds.loc[:10,:]
```

Out[754]:

	cut	color	clarity	depth	table	price	x	y	z
0	5	6	2	0.513889	0.230769	0.000000	0.031384	0.005433	0.044256
1	4	6	3	0.466667	0.346154	0.000000	0.022825	0.002898	0.040351
2	2	6	5	0.386111	0.423077	0.000054	0.045649	0.007063	0.040351
3	4	2	4	0.538889	0.288462	0.000433	0.067047	0.009960	0.050765
4	2	1	2	0.563889	0.288462	0.000487	0.087019	0.012133	0.054670
5	3	1	6	0.550000	0.269231	0.000541	0.029957	0.005071	0.045884
6	3	2	7	0.536111	0.269231	0.000541	0.031384	0.005433	0.045558
7	3	3	3	0.525000	0.230769	0.000595	0.048502	0.007787	0.047511
8	1	6	4	0.613889	0.346154	0.000595	0.019971	0.001811	0.046209
9	3	3	5	0.455556	0.346154	0.000649	0.038516	0.006700	0.042955
10	2	1	3	0.583333	0.230769	0.000703	0.074180	0.010866	0.054019

Test and Training Data

In [755]:

```
train = ds.sample(frac = 0.75, replace = False)
test = ds.drop(train.index)
```

In [756]:

```
train = train.to_numpy()
```

```
train = train.to_numpy(),
test = test.to_numpy()
```

In [757]:

```
label_col = 5
# label_col is price column
y_train = train[:,label_col]
x_train = np.delete(train,label_col,1)
#inserting bias b, this will work for both univariate and multivariate
x_train = np.insert(x_train, 0, np.ones(len(x_train)), axis=1)

y_test = test[:,label_col]
x_test = np.delete(test,label_col,1)
#inserting bias b
x_test = np.insert(x_test, 0, np.ones(len(y_test)), axis=1)
```

In [758]:

```
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(40331, 9)
(13444, 9)
(40331,)
(13444,)
```

Linear Regression

Closed Form

$$\mathbf{w} = (\mathbf{X.TX})^{-1} \mathbf{X.TY}$$

In [759]:

```
def closedForm(x, y):
    return np.dot(np.dot(np.linalg.inv(np.dot(x.T, x)), x.T), y)
```

Gradient Descent

In \mathbf{X} , we are adding a column of all 1's which represents our bias (b) from the equation $y = \mathbf{w} \cdot \mathbf{x} + b$.

$$y_{\text{cap}} = y_{\text{pred}} = \mathbf{x} \cdot \mathbf{w} \quad \text{cost} = J(\mathbf{w}) = -1/m \sum (\mathbf{y} - y_{\text{cap}})^2 \text{ for } i = 1 \text{ to } m$$

$$dJ(\mathbf{w}) / d\mathbf{w} = -1/m \mathbf{X}(\mathbf{y} - y_{\text{cap}}); m = \text{rows}$$

$$a = \text{alpha (learning rate)}$$

$$\mathbf{w} = \mathbf{w} - a * dJ/d\mathbf{w}$$

In [760]:

```
def gradientDescent(x, y, w, itr, alpha):
    costs_list = [1]*itr
    for i in range(itr):
        y_cap = np.dot(x, w)
        loss = y_cap - y
        cost = np.sum(loss ** 2) / len(x)

        #each iterations cost
        costs_list[i] = cost
        if (i % 50 == 0):
            print("Cost=", cost)
            #w = w - alpha * d(loss)/dw
```

```
grad = np.dot(x.T, loss) / len(x)
w = w - alpha * grad
return w, costs_list
```

Mean Squared Error

The Mean Squared Error measures how close a regression line is to a set of data points.

In [761]:

```
def meanSqErr(y_pred, y_test):  
    MeanSquaredError = np.mean(np.square(y_pred - y_test))  
    return MeanSquaredError
```

For Graph Plot

In [762]:

```
def plotCost(cost, method):
    plt.plot(cost)
    plt.ylabel('Cost')
    plt.xlabel('iteration')
    plt.title('Cost curve (' + method + ' solution)')
    plt.show()
```

UniVariate


Lets check the correlation Matrix and check which columns are most dependent on our label column


In [763]:

```
plt.figure(figsize=(11,11))
sns.heatmap(ds.corr(), annot=True)
```

Out[763]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8950f421d0>
```





In [764]:

```
x_train
```

Out[764]:

```
array([[1.          , 5.          , 7.          , ..., 0.06419401, 0.00851141,
        0.04913765],
       [1.          , 3.          , 7.          , ..., 0.2810271 , 0.03730532,
        0.07972665],
       [1.          , 4.          , 3.          , ..., 0.53780314, 0.06700471,
        0.11096648],
       ...,
       [1.          , 5.          , 4.          , ..., 0.22253923, 0.02951829,
        0.07159128],
       [1.          , 4.          , 5.          , ..., 0.41940086, 0.05378486,
        0.10055321],
       [1.          , 2.          , 1.          , ..., 0.50071327, 0.06519377,
        0.11519688]])
```

In [765]:

```
#here we are taking two columns from our x_train and x_test
# 1st column = bias (b) and i'th column = feature on which we will apply Univariate.
def train_test_univariate(i):
    #                bias                i'th column
    ux_train = np.array([x_train[:,0], x_train[:,i]]).T
    ux_test = np.array([x_test[:,0], x_test[:,i]]).T
    return ux_train, ux_test, y_train, y_test
```

We can take any of the columns : x, y, z for prediction of label

In [766]:

```
x_col = ds.columns.get_loc("x")
y_col = ds.columns.get_loc("y")
z_col = ds.columns.get_loc("z")
```

For X Column

In [767]:

```
ux_train, ux_test, uy_train, uy_test = train_test_univariate(x_col)
uy_train=uy_train.reshape((len(uy_train), 1))
uy_test=uy_test.reshape((len(uy_test), 1))
```

In [768]:

```
#Closed Form Solutin

weightC_UNI = closedForm(ux_train, uy_train)
y_pred_closed_UNI = np.dot(ux_test, weightC_UNI)
print("Closed form UniVariate using << X >> column Mean Squared Error = ", meanSqErr(y_pred_closed_UNI, uy_test))
```

Closed form UniVariate using << X >> column Mean Squared Error = 0.010007017074549722

In [769]:

```
#Gradient Descent Solution
```

```
weightG_UNI = 0.1
itr = 600
alpha = 0.35
weightG_UNI, costsG_UNI = gradientDescent(ux_train, uy_train, weightG_UNI, itr, alpha)
```

```
Cost= 0.1230726656888536
Cost= 0.04556362687616016
Cost= 0.03102156416809773
Cost= 0.024656724709065096
Cost= 0.02187093137612901
Cost= 0.020651632045205953
Cost= 0.02011796329066754
Cost= 0.019884384601376395
Cost= 0.019782150765074398
Cost= 0.019737404571332398
Cost= 0.019717819843547355
Cost= 0.019709247906353627
```


In [770]:

```
y_pred_grad_UNI = np.dot(ux_test, weightG_UNI)
print("Gradient Descent UniVariate using <X> column Mean Squared Error = ", meanSqErr(y_pred_grad_UNI, uy_test))
```

Gradient Descent UniVariate using <X> column Mean Squared Error = 0.010011437186306944

In [771]:

```
plotCost(costsG_UNI, "Gradient descent using UniVariate: ")
```



For Y Column

In [772]:

```
ux_train, ux_test, uy_train, uy_test = train_test_univariate(y_col)
uy_train=uy_train.reshape((len(uy_train), 1))
uy_test=uy_test.reshape((len(uy_test), 1))
```

In [773]:

```
#Closed Form Solution
```

```
weightC_UNI = closedForm(ux_train, uy_train)
y_pred_closed_UNI = np.dot(ux_test, weightC_UNI)
print("Closed form UniVariate using << Y >> column Mean Squared Error = ", meanSqErr(y_pred_closed_UNI, uy_test))
```

Closed form UniVariate using << Y >> column Mean Squared Error = 0.017043747722067813

In [774]:

```
#Gradient Descent Solution
weightG_UNI = 0.1
```

```
itr = 1000
alpha = 0.011
weightG_UNI, costsG_UNI = gradientDescent(ux_train, uy_train, weightG_UNI, itr, alpha)
```

```
Cost= 0.13697936620553774
Cost= 0.09631344469359934
Cost= 0.09269088319243472
Cost= 0.09147294628003394
Cost= 0.09104936538727677
Cost= 0.09088820245596667
Cost= 0.09081373642459524
Cost= 0.09076791938236774
Cost= 0.0907315757084925
Cost= 0.09069837090116778
Cost= 0.09066621240927056
Cost= 0.09063440897276392
Cost= 0.090602732235794
Cost= 0.0905711067588329
Cost= 0.09053950761691729
Cost= 0.09050792657212775
Cost= 0.09047636089896492
Cost= 0.09044480969291056
Cost= 0.09041327265099361
Cost= 0.09038174966896048
```


In [775]:

```
y_pred_grad_UNI = np.dot(ux_test, weightG_UNI)
print("Gradient Descent UniVariate using << Y >> column Mean Squared Error = ", meanSqErr
(y_pred_grad_UNI, uy_test))
```

```
Gradient Descent UniVariate using << Y >> column Mean Squared Error = 0.0457747901126378
3
```

In [776]:

```
plotCost(costsG_UNI, "Gradient descent using UniVariate: ")
```



Z Column

In [777]:

```
#Gradient Descent Solution
weightG_UNI = 0.1
itr = 1000
alpha = 0.011
weightG_UNI, costsG_UNI = gradientDescent(ux_train, uy_train, weightG_UNI, itr, alpha)
```

```
Cost= 0.13697936620553774
Cost= 0.09631344469359934
Cost= 0.09269088319243472
Cost= 0.09147294628003394
Cost= 0.09104936538727677
Cost= 0.09088820245596667
```

```
Cost= 0.09000020243330007
Cost= 0.09081373642459524
Cost= 0.09076791938236774
Cost= 0.0907315757084925
Cost= 0.09069837090116778
Cost= 0.09066621240927056
Cost= 0.09063440897276392
Cost= 0.090602732235794
Cost= 0.0905711067588329
Cost= 0.09053950761691729
Cost= 0.09050792657212775
Cost= 0.09047636089896492
Cost= 0.09044480969291056
Cost= 0.09041327265099361
Cost= 0.09038174966896048
```

In [778]:

```
ux_train, ux_test, uy_train, uy_test = train_test_univariate(z_col)
uy_train=uy_train.reshape((len(uy_train), 1))
uy_test=uy_test.reshape((len(uy_test), 1))
```

In [779]:

```
#Closed Form Solutin
```

```
weightC_UNI = closedForm(ux_train, uy_train)
y_pred_closed_UNI = np.dot(ux_test, weightC_UNI)
print("Closed form UniVariate using << Z >> column Mean Squared Error = ", meanSqErr(y_pr
ed_closed_UNI, uy_test))
```

Closed form UniVariate using << Z >> column Mean Squared Error = 0.014940489937509394

In [780]:

```
#Gradient Descent Solution
```

```
weightG_UNI = 0.1
itr = 1000
alpha = 0.012
weightG_UNI, costsG_UNI = gradientDescent(ux_train, uy_train, weightG_UNI, itr, alpha)
```

```
Cost= 0.1352712007544064
Cost= 0.09508750523039379
Cost= 0.09196686550196595
Cost= 0.09101135470391955
Cost= 0.09069821895911043
Cost= 0.09057568290888587
Cost= 0.09050971242647102
Cost= 0.09046054184024391
Cost= 0.09041637333653367
Cost= 0.09037370669260293
Cost= 0.09033150344076013
Cost= 0.09028945547445004
Cost= 0.09024747137293321
Cost= 0.09020552400249734
Cost= 0.09016360530231136
Cost= 0.09012171287015225
Cost= 0.0900798459826169
Cost= 0.09003800441438607
Cost= 0.0899961880879201
Cost= 0.0899543969695295
```

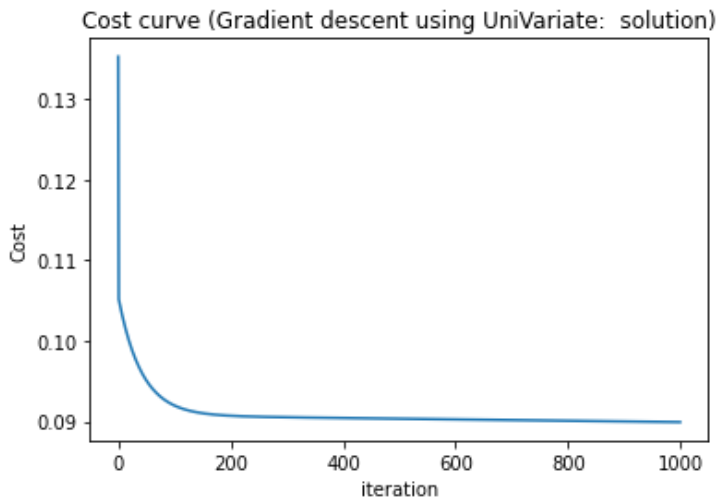
In [781]:

```
y_pred_grad_UNI = np.dot(ux_test, weightG_UNI)
print("Gradient Descent UniVariate using << Z >> column Mean Squared Error = ", meanSqErr
(y_pred_grad_UNI, uy_test))
```

Gradient Descent UniVariate using << Z >> column Mean Squared Error = 0.0455579780196802
2

In [782]:

```
plotCost(costsG_UNI, "Gradient descent using UniVariate: ")
```



MultiVariate

In [783]:

```
#Closed Form Solution
```

```
weightC = closedForm(x_train, y_train)
```

```
#Gradient Descent Solution
```

```
weightG = [0.1] * len(x_train[0])
```

```
itr = 220
```

```
alpha = 0.001
```

```
weightG, costsG = gradientDescent(x_train, y_train, weightG, itr, alpha)
```

```
Cost= 1.7276138171175495
```

```
Cost= 0.04461621990072215
```

```
Cost= 0.038621204751832144
```

```
Cost= 0.038480786396332324
```

```
Cost= 0.03836601866517273
```

In [784]:

```
#closed form prediction
```

```
y_pred_closed = np.dot(x_test, weightC)
```

```
#gradient descent prediction
```

```
y_pred_grad = np.dot(x_test, weightG)
```

In [785]:

```
print("Closed form Mean Squared Error = ", meanSqErr(y_pred_closed, y_test))
```

```
Closed form Mean Squared Error = 0.011663738954442171
```


In [786]:


```
print("Gradient descent Mean Squared Error = ", meanSqErr(y_pred_grad, y_test))
```

```
Gradient descent Mean Squared Error = 0.03892606575592581
```

In [787]:

```
plotCost(costsG, "Gradient descent")
```






Further checking for correlated columns and checking for results

In [788]:

```
plt.figure(figsize=(11,11))
sns.heatmap(ds.corr(), annot=True)
```

Out[788]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f895031fe50>



Let's also try with dropping x, y and z values as they are highly correlated to our target column

In [789]:

```
x_col = ds.columns.get_loc("x")
y_col = ds.columns.get_loc("y")
z_col = ds.columns.get_loc("z")
```

In [790]:

```
ds.head()
```

Out[790]:

	cut	color	clarity	depth	table	price	x	y	z
0	5	6	2	0.513889	0.230769	0.000000	0.031384	0.005433	0.044256
1	4	6	3	0.466667	0.346154	0.000000	0.022825	0.002898	0.040351
2	2	6	5	0.386111	0.423077	0.000054	0.045649	0.007063	0.040351
3	4	2	4	0.538889	0.288462	0.000433	0.067047	0.009960	0.050765
4	2	1	2	0.563889	0.288462	0.000487	0.087019	0.012133	0.054670

In [791]:

```
ds.drop(axis="columns", labels="y", inplace=True)
ds.drop(axis="columns", labels="x", inplace=True)
ds.drop(axis="columns", labels="z", inplace=True)
```

Again calculating x_train, x_test, y_train and y_test after dropping columns and checking for multivariate values

In [792]:

```
ds.head()
```

Out[792]:

	cut	color	clarity	depth	table	price
0	5	6	2	0.513889	0.230769	0.000000
1	4	6	3	0.466667	0.346154	0.000000
2	2	6	5	0.386111	0.423077	0.000054
3	4	2	4	0.538889	0.288462	0.000433
4	2	1	2	0.563889	0.288462	0.000487

Removing from test and train sets

In [793]:

```
label_col = 5
# label_col is price column
y_train = train[:,label_col]

#removing label_column and x, y, z columns
x_train = np.delete(train,[label_col, x_col, y_col, z_col],1)

#inserting bias b, this will work for both univariate and multivariate
x_train = np.insert(x_train, 0, np.ones(len(x_train)), axis=1)

y_test = test[:,label_col]

#removing label_column and x, y, z columns
x_test = np.delete(test,[label_col, x_col, y_col, z_col],1)

#inserting bias b
x_test = np.insert(x_test, 0, np.ones(len(y_test)), axis=1)
```

Final Solution after removing all the correlated columns

In [794]:

```
#Closed Form Solution
```

```
weightC = closedForm(x_train, y_train)
```

```
#Gradient Descent Solution
```

```
weightG = [0.1] * len(x_train[0])
```

```
itr = 220
```

```
alpha = 0.001
```

```
weightG, costsG = gradientDescent(x_train, y_train, weightG, itr, alpha)
```

```
Cost= 1.6398328097028203
```

```
Cost= 0.05199669078742636
```

```
Cost= 0.046331486882540414
```

```
Cost= 0.046238802235241565
```

```
Cost= 0.04617420689155505
```

```
In [795]:
```

```
#closed form prediction
```

```
y_pred_closed = np.dot(x_test, weightC)
```

```
#gradient descent prediction
```

```
y_pred_grad = np.dot(x_test, weightG)
```

```
In [796]:
```

```
print("Closed form Mean Squared Error = ", meanSqErr(y_pred_closed, y_test))
```

```
Closed form Mean Squared Error = 0.044246800117223595
```

```
In [797]:
```

```
print("Gradient descent Mean Squared Error = ", meanSqErr(y_pred_grad, y_test))
```

```
Gradient descent Mean Squared Error = 0.04680266014458536
```

```
In [798]:
```

```
plotCost(costsG, "Gradient descent")
```

